A Survey of Dimensionality Reduction Techniques

RACHEL LEVANGER

This note set gives a short summary of a collection of commonly-used linear and non-linear dimensionality reduction techniques.

Contents

1   Linear Dimensionality Reduction (LDR) Techniques

1.1   Principal Component Analysis (PCA)

PCA is a tool used to describe a data set in terms of its extrinsic (linear) directions of highest variance, and does so in a way that the data is uncorrelated with respect to these directions. Consider a set S of m points in $\mathbb{R}^n$. Then S can be represented as an $m \times n$ matrix X. Call $\mu_X = (\mu_1, ..., \mu_n)$ the center of the data set, or the average of each of the columns of X. The general sequence of steps to perform a PCA on X is as follows:

(1)   Traditionally, compute the covariance matrix Σ associated to the data set X, where

$$\Sigma_{ij} = \mathrm{cov}(X_{*i}, X_{*j}) = (1/m) \sum_{k=1}^{m} (X_{ki} - \mu_i)(X_{kj} - \mu_j).$$

Notice that this is just the dot product of the mean-centered column vectors of the data matrix divided by the number of points in S. It is also possible to use $\Sigma = X^T X$ or the correlation matrix, as these are all essentially functions of the dot product, they are real and symmetric, and so diagonalization yields an orthogonal set of eigenvectors. Using $\Sigma = X^T X$ is essentially the computation of the Singular Value Decomposition (SVD) of X as it relates to finding the directions of highest variance.

(2)   Since Σ is real and symmetric, we can diagonalize Σ via the matrices Φ and Λ, where Φ is orthogonal and Λ is diagonal. The matrix Φ has columns that form the orthonormal basis of eigenvectors of Σ, Λ gives the corresponding eigenvalues, and the equation $\Sigma = \Phi \Lambda \Phi^T$ holds.

(3)   The magnitude of the eigenvalues in Λ provide insight into the number of dimensions along which the data is organized. Furthermore, since the eigenvectors are orthogonal, the covariance along any two transformed basis vectors are zero (their dot product is zero).

sources:
https://stat.duke.edu/ sayan/Sta613/2015/lec/IDAPILecture15.pdf
http://infolab.stanford.edu/ ullman/mmds/ch11.pdf

1.2 Multidimensional scaling (MDS)

2 Non-linear Dimensionality Reduction (NLDR) Techniques

2.1 Kernel PCA

2.2 Isomap

2.3 Locally-linear Embedding (LLE)

This technique works best when the points in your dataset are assumed to have densely sampled a portion of a manifold that can be 'flattened out' in some sense. For instance, the swiss roll is a two-dimensional manifold with boundary that can be flattened out into a portion of a two-dimensional plane, while a torus cannot be flattened out as such and would not be a good candidate for this method.

Locally-linear embedding (LLE) rests on the assumption that points that are nearest to a target point lie on the same portion of the manifold as the target point, i.e. the sampling density is much smaller than the weak feature size of the underlying manifold. Let $X = \{x_1, ..., x_N\}$ be a set of points in $\mathbb{R}^D$, $K \in \mathbb{N}$, and $d < D$ be the desired embedding dimension.

Algorithm:

(1) *Find the K nearest neighbors of each point.* Compute the distance matrix of the data set X. Let $N_K : X \to \mathcal{P}(X)$ be the function that assigns to each x_i its K nearest neighbors, $N_K(x_i)$, which can be gleaned from the distance matrix.

(2) *Solve for the reconstruction weights.*

    (a) For each $x_i \in X$, create a matrix Z_i with a column for each element in $N_K(x_i)$. Subtract x_i from each column in Z (center at x_i).

    (b) Set $C_i = Z_i^T Z_i$ (K times the local covariance), and solve for a vector w_i such that $C_i w_i = [1]$, where $[1]$ is a column vector of all ones.

    (c) Populate the i^{th} row of the weight matrix W as follows. If $x_j \notin N_K(x_i)$, set $W_{ij} = 0$. If $x_j \in N_K(x_i)$ then set $W_{ij} = w_j/(w \cdot [1])$, where w_j is the entry in w corresponding to the neighbor x_j.

(3) *Compute the embedding coordinates.* Set $M = (I - W)^T (I - W)$. Find the $d + 1$ eigenvectors corresponding to the smallest eigenvalues of M. Set the q^{th} row of Y to be the eigenvector corresponding to the $(q + 1)^{th}$-smallest eigenvalue.

(Thus, the eigenvector corresponding to the smallest eigenvalue of M is not used in Y.

sources:
https://www.cs.nyu.edu/ roweis/lle/algorithm.html

2.4 Hessian Locally-linear Embedding (Hessian LLE)

2.5 Laplacian Eigenmaps

2.6 Diffusion Maps

rachel@math.rutgers.edu